香港中文大學
The Chinese University of Hong Kong

Reinforcement Learning: function approximation

Zihao DENG
Wenqian ZHAO

Department of Computer Science and Engineering
The Chinese University of Hong Kong

Markov Decision Process

- The MDP gives us a precise formulation of the environment, given a state st we select an action at and observe s_{t+1} and r_t according to the transition probabilities P:
  - S - set of possible states
  - $s_t \in S$ - state at step t
  - A - set of possible actions
  - $a_t \in A$ - selected action at step t
  - R - Reward function. The reward at step t is given by $r_{t+1} = R(s_t, a_t, s_{t+1})$
  - P - transition probabilities such that $s_{t+1} P(s|s_t, a_t)$, i.e.
  - ρ - Initial state distribution such that $s_0 \rho(s)$
- Agent is defined with a policy function $\pi(a|s)$, mapping from states to actions and can be either deterministic or non-deterministic

Markov Decision Process

- Given an MDP and a policy, an episode can be produced by repeating of:
  - $a_t \pi(a|s_t)$
  - $s_{t+1} P(s|s_t, a_t)$
  - $r_{t+1} = r(s_t, a_t, s_{t+1})$
- which produce:

$$\text{episode} := s_0, a_0, r_1, s_1, a_1, r_1, \ldots, s_{\mathcal{T}1}, a_{\mathcal{T}1}, r_{\mathcal{T}1}, s_{\mathcal{T}}$$

- Optimal Solution gives:

$$\max_{\pi} E[\sum_{t=1}^{\mathcal{T}} r_t] \tag{1}$$

Markov Decision Process

- Value function defined as :

$$V^{\pi}(s) = E_{\pi}[\sum_{t=1}^{\tau} r_t | s_0 = s] \tag{2}$$

$$V^{*}(s) = \max_{\pi} E_{\pi}[\sum_{t=1}^{\tau} r_t | s_0 = s] \tag{3}$$

- Bellman equation: A recursive relation for value function:

$$V^{*}(s) = \max_{a \in A} E[r_{t+1} + V^{*}_{s_{t+1}} | s_t = s, a_t = t] \tag{4}$$

- $(TV)(s)$ is the Bellman operator and we can recursively calculate it and update $V(s)$ (value iteration) to reach optimal (Monotonicity and Contraction mapping)

$$(TV)(s) = \max_{a \in A} E[r_{t+1} + V^{*}_{s_{t+1}} | s_t = s, a_t = t] \tag{5}$$

$$V_{k+1} = TV_k \tag{6}$$

- policy iteration use the same idea, but instead of updating value function, it update Policy π

Markov Decision Process

- Another approach: State-Value Function: define a quantity $Q : S \times A \to \mathbb{R}$:

$$Q^{\pi}(s,a) = \bar{R}(s,a) + \sum_{s' \in S}^{T} P_{s,a}(s') V^{\pi}(s') \tag{7}$$

- Recursively Calculate optimal by using Bellman Operator:
  - $FQ(s,a) = \bar{R}(s,a) + \sum_{s' \in S}^{T} P_{s,a}(s') \max_{a' \in A} Q(s', a')$
  - $Q(s,a) = FQ(s,a)$

- Greedy action selection is simple:

$$\pi(s) = \arg\max_{a' \in A} Q(s_{t+1}, a') \tag{8}$$

MDP vs RL

- Difference between Markov Decision Process and Reinforcement learning:

  - MDP: the transition matrix $P(s'|s,a)$ is known $\rightarrow$ used to find the optimal agent
  - RL: $P(s'|s,a)$ unknown and need to be learned :
    i. Iteracting with environment
    ii. Requiring explicit knowledge of P

- But the same idea of state-action value function can be used for Reinforcement learning

Q-learning and SARSA

- Q-learning:

  i. Build a Q-table which stores $Q(s, a)$ for each s and a (randomly initialized). i.e.

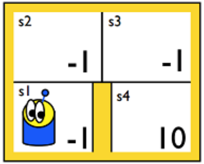

$\alpha = .7$

	↑	↓	←	→
s_1	0	0	0	0
s_2	0	0	0	0
s_3	0	0	0	0
s_4	0	0	0	0

Q-Table

Q-learning and SARSA

- Q-learning:

  ii. update $Q(s, a)$ with:

$$Q_{k+1} := (1 - \gamma_k)Q_k + \gamma_k(r + \max_{a' \in A} Q_k(s', a')) \tag{9}$$

  where γ_k is the learning rate, with $\sum_{k=1}^{\infty} \gamma_k = \infty$ and $\sum_{k=1}^{\infty} \gamma_k^2 < \infty$:

$$Q_{k+1} = Q_k + \gamma_k(r + \max_{a' \in A} Q_k(s', a') - Q_k(s, a)) \tag{10}$$

- Q-learning Demo: google

Q-learning and SARSA

- Q-learning result:

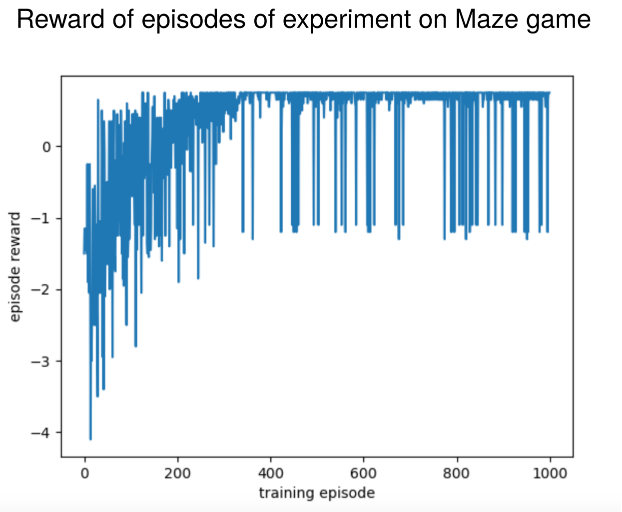

Q-learning and SARSA

- Exploration: $\epsilon - greedy$

$$a_t = \begin{cases} \arg\max_{a \in A} Q_{s_t,a}, & \text{w.p. } 1-\epsilon. \\ unif(A), & \text{w.p. } \epsilon. \end{cases} \tag{11}$$

- SARSA:
  - update based on the current play (s, a, r, s', a')

$$Q_{k+1} = Q_k + \gamma_k(r + Q_k(s', a') - Q_k(s, a)) \tag{12}$$

  - Similar to Q-learning but is On-policy

Deep Q-Networks

- Drawback of Q-learning and SARSA:
  - Q-table can be too big if environment is complicate i.e. $1^6 \times 1^3$ maze
- Alternative Algorithm: DQN:
  - Use a function approximator to estimate action-value function with Q-Network
- Steps:
  - store transition$(s_t, s_t, r_{t+1}, s_{t+1})$ in memory
  - sample mini-batch of transitions, optimise MSE between Q-network and Q-learning targets:

$$\text{minimize } L_w = E_{s,a,r,s'}[(r + \gamma \max_{a'} Q(s', a'; w^-) - Q(s, a; w))^2] \tag{13}$$

  - Important tricks: experience replay and fixed target

Deep Q-Networks(DQN): play games in OpenAI gym

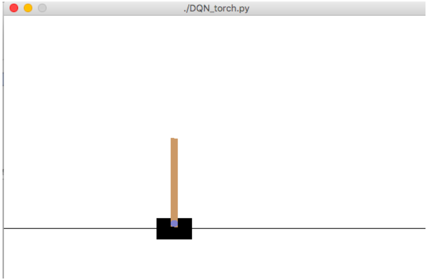

- States are represented by 4-element tuples (position,cart velocity,angle,tip velocity)
- Actions can be either moving left or right
- Function approximator is a feed foward neural network
- 1 hidden layer with 10 neurons, 2 output neurons representing value estimation for two actions
- Implemented using torch and tensorflow, can stay alive for 1 minute

Policy based method

What's wrong with value based methods?

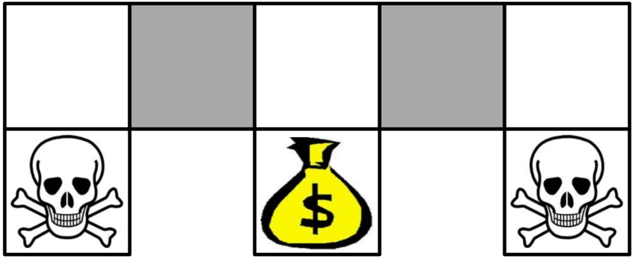

- Main problem: deterministic policy
- No good for partially observable environment
- The agent cannot differentiate the grey states (Horizontally symmetric)
- An optimal deterministic will either go left or right
- In this case, if change starting point, the agent may not reach optimal:run

Policy Gradient: problem formulation

- Policy is a function of observation: $\pi_\theta(\cdot)$
- Trajectory τ: $\{s_0, a_0, r_0, s_1, a_1, r_1, ...\}$ is treated as random variable
- Distribution of τ is determined by policy π_θ
- For each trajectory, total reward is defined as $R(\tau)$
- Ultimate goal: optimize expectation $E_{\pi_\theta}[R(\tau)]$ w.r.t θ

Policy Gradient: approximate the gradient

- What does the gradient look like?

$$
\begin{aligned}
\nabla_\theta E_{\pi_\theta}[R(\tau)] &= \nabla_\theta \sum_\tau P_\theta(\tau)R(\tau) \\
&= \sum_\tau \nabla_\theta P_\theta(\tau)R(\tau) \\
&= \sum_\tau P_\theta(\tau)\frac{\nabla_\theta P_\theta(\tau)}{P_\theta(\tau)}R(\tau) = E_{\pi_\theta}[\nabla_\theta \ln P_\theta(\tau)R(\tau)]
\end{aligned}
\tag{14}
$$

- Equation 14 tell us: gradient can be represented as an expectation
- Why it is important: expectation can be approximated by sampling

Policy Gradient: approximate the gradient

- Why the gradient even exists?

$$P(\tau) = P(s_0)\prod_{i=0}^{\infty}\pi_\theta(a_i, s_i)P(s_{i+1}|s_i, a_i) \tag{15}$$

- Assumption: there is an underlying MDP specifying $P(s_{i+1}|s_i, a_i)$ and $P(s_0)$

$$\begin{aligned}\nabla_\theta \ln P(\tau) &= \nabla_\theta \ln[P(s_0)\prod_{i=0}^{\infty}\pi_\theta(a_i, s_i)P(s_{i+1}|s_i, a_i)] \\ &= \nabla_\theta \ln P(s_0) + \nabla_\theta \sum_{i=0}^{\infty}[\ln \pi_\theta(a_i, s_i) + \ln P(s_{i+1}|s_i, a_i)] \\ &= \nabla_\theta \sum_{i=0}^{\infty} \ln \pi_\theta(a_i, s_i)\end{aligned} \tag{16}$$

Policy Gradient: understanding the formula

Combine all equation in previous slides, one important formula:

$$\nabla_\theta E_{\pi_\theta}[R(\tau)] = E_{\pi_\theta}[R(\tau)\nabla_\theta \sum_{s_i,a_i \in \tau} \ln \pi_\theta(a_i, s_i)] \tag{17}$$

Intuition from equation 17, adjustment magnitude of policy on $\pi_\theta(a, s)$:

- In proportion to the total reward gained from trajectories containing (a, s)
  - Rationale: good actions lead to good trajectories, while bad actions lead to bad ones
  - What about good actions in trajectories with bad overall performance? work on it later
- In inverse proportion to the probability of performing action a on state s
  - consider actions sampled frequently but with small positive effect
  - mitigate case where 'not-so-good' actions are rewarded frequently

Vanilla Policy Gradient: **REINFORCE**

So far, we obtain the first policy gradient algotithm called **REINFORCE** [Williams, R. J.]

Experiment on CartPole game:

Bad performance, worse than random play after 1000 episodes of training

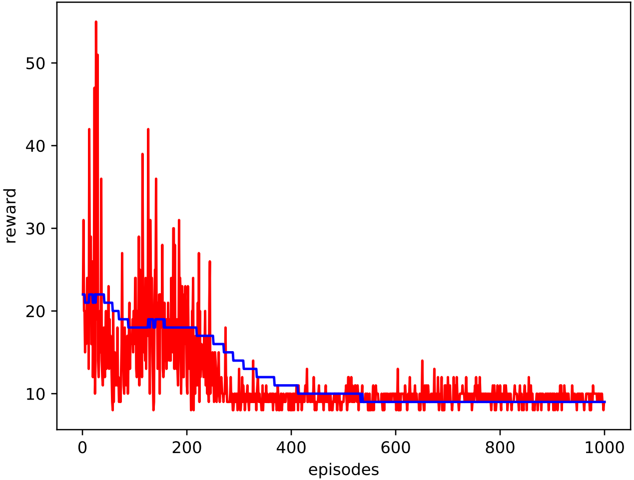

Improvement for **REINFORCE**: baseline

- One natural quetion: what if reward is always positive?
- Do we have to always increase $\pi_\theta(a, s)$ because $R(\tau)$ is positive (as in formula 17)?
- Actually, we only cares about the relative performance of trajectories
- Observation from formula 18: we can remove any constant term A from the expectation with out introducing bias.
- A can be the average performance for all trajectories, it is referenced as a baseline

$$\begin{aligned} E_{\pi_\theta}[\sum_a A \cdot \nabla \ln \pi(a, S)] &= \sum_a \pi_\theta(a, S) A \frac{\nabla_\theta \pi_\theta(a, S)}{\pi_\theta(a, S)} \\ &= A \sum_a \nabla_\theta \pi(a, S) \\ &= A \cdot \nabla_\theta \sum_a \pi_\theta(a, S) = A \nabla_\theta(1) = 0 \end{aligned} \tag{18}$$

Improvement for **REINFORCE**: baseline

Experiment on CartPole game:
Better than before: an upgoing trend of rewards

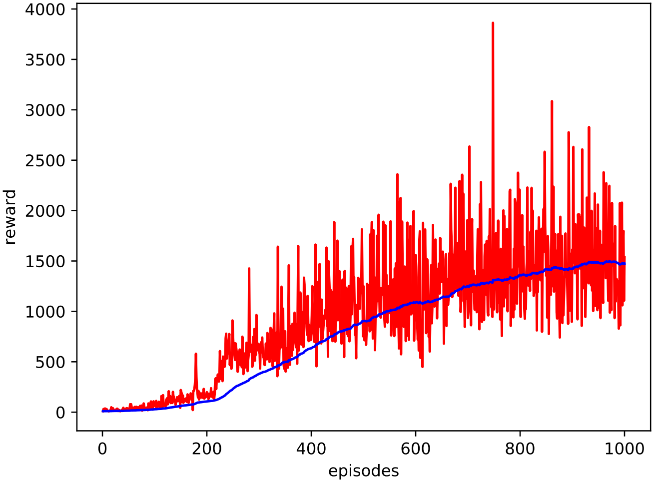

Improvement for **REINFORCE**: advantage function

- Why we award/punish an action (s,a) based on the entire trajectory reward?
- Markov property: the action a_t only affects rewards after time t.

$$E_{\pi_\theta}[R_{0:i-1}(\tau)\nabla_\theta \ln \pi_\theta(a_i, s_i)] = 0 \tag{19}$$

- Actually, we can exploit the markov property to refine formula 17

$$\begin{aligned}\nabla_\theta E_{\pi_\theta}[R(\tau)] &= E_{\pi_\theta}[\sum(\nabla_\theta \ln \pi_\theta(a_i, s_i)(R_{0:i-1}(\tau) + R_{i:\infty}(\tau)))] \\ &= E_{\pi_\theta}[\sum(\nabla_\theta \ln \pi_\theta(a_i, s_i)R_{i:\infty}(\tau))]\end{aligned} \tag{20}$$

- Recall that subtraction of baseline doesn't change the expectation

$$\nabla_\theta E_{\pi_\theta}[R(\tau)] = E_{\pi_\theta}[\sum \nabla_\theta \ln \pi_\theta(a_i, s_i)(R_{i:\infty}(\tau) - V_{\pi_\theta}(s_i))] \tag{21}$$

Improvement for **REINFORCE**: advantage function

- As shown in equation 21, $R_{i:\infty}(\tau) - V_{\pi_\theta}(s_i)$ is the actual term that determines the magnitude we adjust probability $\pi_\theta(a_i, s_i)$
- $R_{i:\infty}(\tau) - V_{\pi_\theta}(s_i)$ is also known as the advantage function
- Rationale: extra reward gained when performing certain action a_i on state s_i compared to average reward from that state under policy π_θ

Improvement for **REINFORCE**: advantage function

Experiment on CartPole game:

Only after 141 episodes of training, surviving time boosted to 20k!

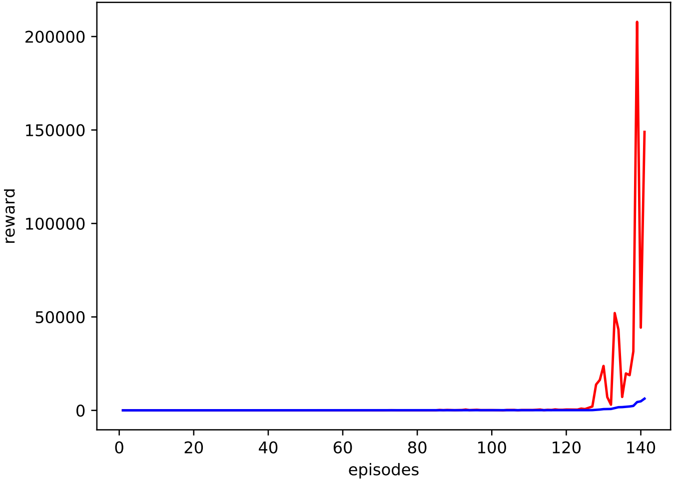

Actor-Critic: a combination

So far, we mainly focused on pure value-based and pure policy-based methods ...

- Value-based: problem of deterministic policy in partially observed environments
- Policy-based: credit assignment problem (delay between action and reward)
- Why not combine them?
- Still use policy function
- Also adopt an estimator for state values to approximate advantage function
- Policy updates without delay!

Further improvement: clipped objective function

- In previous section, policy gradient method works by computing gradient estimator in form

$$\hat{g} = \hat{E}_t[\nabla_\theta \ln \pi_\theta(a_t, s_t)\hat{A}_t] \tag{22}$$

- The estimator $\hat{g}$ can be obtained by differentiating the objective

$$L^{PG}(\theta) = \hat{E}_t[\ln \pi_\theta(a_t, s_t)\hat{A}_t] \tag{23}$$

- Multiple steps to optimize L^{PG} on same trajectory: destructively large policy updates [Schulman, John et al.].
- To improve sample efficiency, they adopt strategy of clipping the surrogate objective function in form $L^{CPI}(\theta)$:

Further improvement: clipped objective function

$$L^{CPI}(\theta) = \hat{E}_t[\frac{\pi_\theta(a_t, s_t)}{\pi_{\theta_{old}}(a_t, s_t)}\hat{A}_t] \tag{24}$$

- $\pi_{\theta_{old}}$: fixed term generated by old policy
- π_θ: current policy being optimized.
- The ratio $\frac{\pi_\theta(a_t,s_t)}{\pi_{\theta_{old}}(a_t,s_t)}$ is denoted as $r_t(\theta)$
- $r_t(\theta)$ measures the difference between current policy and old policy
- we don't want too big a update step, hence some constraint based on $r_t(\theta)$
- In practise we use the gradient of following objective function

$$L^{CLIP}(\theta) = \hat{E}_t[\min(r_t(\theta)\hat{A}_t, \text{clip}(r_t(\theta), 1-\epsilon, 1+\epsilon)\hat{A}_t] \tag{25}$$

Further improvement: clipped objective function

- Why $\hat{E}_t[\frac{\pi_\theta(a_t,s_t)}{\pi_{\theta_{old}}(a_t,s_t)}\hat{A}_t]$ in equation 24?
- Well justified in "Trust Region Policy Optimization"[Schulman, John et al.]
- My understanding: a case of importance sampling
- Importance sampling: adjusted rewards, learn from different policy
- Trajectories generated from π_{old} are learned multiple times to update a different policy π, through importance sampling
- More on importance sampling: On a Connection between Importance Sampling and the Likelihood Ratio Policy Gradient [Tang Jie and Pieter Abbeel.]

Proximal Policy Optimization(PPO): some demo

Test on OpenAI gym Agents implemented and trained using Pytorch

For detailed information about task environment, check this list

- CartPole-v0: no training and trained
- MountainCar-v0: no training and trained
- LunarLander-v2: no training and trained
- Pendulum-v0: no training and trained

Some strategyies in our training:

- For continous action space (like Pendulum-v0): discretize it
- Set a maximum number (5000) of steps for each episode during training
- Use a large batch size (512) to perform gradient descent
- Adopt different step size for Actor and Critic updates
- Have a look at our code on github

Further improvement: high score buffer replay

- The learning curve is like:

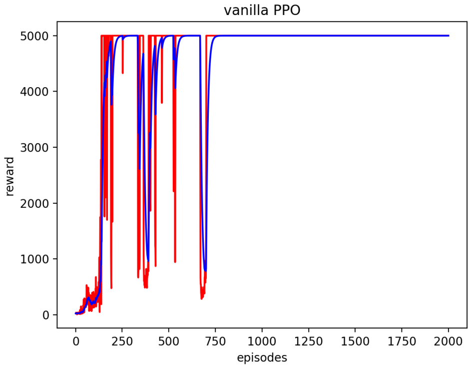

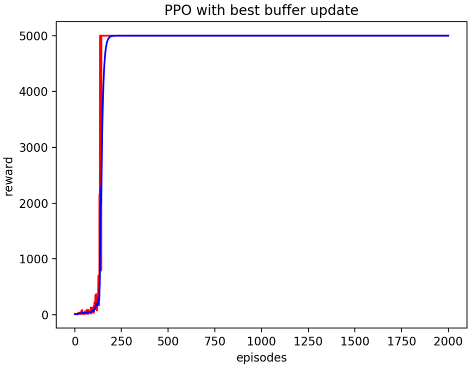

Reference

- *Williams, Ronald J. Simple statistical gradient-following algorithms for connectionist reinforcement learning. Machine learning, 8(3-4):229-256, 1992.*
- *Schulman, J., Levine, S., Moritz, P., Jordan, M.I., & Abbeel, P. (2015). Trust Region Policy Optimization. ICML.*
- *Tang, J., & Abbeel, P. (2010). On a Connection between Importance Sampling and the Likelihood Ratio Policy Gradient. NIPS.*
- *Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). Proximal Policy Optimization Algorithms. CoRR, abs/1707.06347.*

Github link for the whole project: https://github.com/JamesTuna/RL_collects