



香港中文大學  
The Chinese University of Hong Kong


# Reinforcement Learning: function approximation

Zihao DENG  
Wenqian ZHAO

Department of Computer Science and Engineering  
The Chinese University of Hong Kong



# Value Function Approximation



- ▶ Tablular methods: impossible to record all states for real word problems
- ▶ Function approximation: generalize from seen states to unseen states



# Value Function Approximation

- ▶ Goal: find parameter vector  $w$  minimising mean-squared error between approximate value function  $\hat{v}(S, w)$  and true value function  $v_\pi(S)$

$$J(w) = ||v_\pi(S) - \hat{v}(S, w)||_2^2 \quad (1)$$

- ▶ Stochastic gradient descent samples the gradient

$$\Delta w = \alpha(v_\pi(s) - \hat{v}(S, w))\nabla_w \hat{v}(S, w) \quad (2)$$

- ▶ In reality we don't have the true value function  $v_\pi(S)$ 
  - For Monte-Carlo, use discounted return  $G_t$
  - For TD, use  $R_{t+1} + \lambda \hat{v}(S_{t+1}, w)$



# Deep Q-Networks (DQN)


- ▶ Take action  $a_t$  according to  $\epsilon$ -greedy policy
- ▶ Store transition  $(s_t, a_t, r_{t+1}, s_{t+1})$  in memory  $D$
- ▶ Sample random mini-batch of transitions  $(s, a, r, s')$  from  $D$
- ▶ Compute Q-learning targets w.r.t. old, fixed parameters  $w^-$
- ▶ Optimise MSE between Q-network and Q-learning targets

$$L(w) = \mathbb{E}_{s,a,s',r' \sim D_i} [(r + \gamma \max_{a'} Q(s', a'; w^-) - Q(s, a; w))^2] \quad (3)$$

- ▶ Two important tricks in ensuring convergence: experience replay and fixed target




# Deep Q-Networks(DQN): play games in OpenAI gym



- ▶ States are represented by 4-element tuples (position, cart velocity, angle, tip velocity)
- ▶ Actions can be either moving left or right
- ▶ Function approximator is a feed forward neural network
- ▶ 1 hidden layer with 10 neurons, 2 output neurons representing value estimation for two actions
- ▶ Implemented using torch and tensorflow, can stay alive for 1 minute



# Deep Q-Networks(DQN): play games in OpenAI gym



- ▶ States are represented by 4-element tuples (position, cart velocity, angle, tip velocity)
- ▶ Actions can be either moving left or right
- ▶ Function approximator is a feed forward neural network
- ▶ 1 hidden layer with 10 neurons, 2 output neurons representing value estimation for two actions
- ▶ Implemented using torch and tensorflow, can stay alive for 1 minute

